

Enterprise AI T&E Implementation Plan

Turning the AI Test, Evaluation & Assurance Strategy into Executed Work

August 2026

Purpose

The companion *Enterprise AI Test, Evaluation & Assurance Strategy* sets direction: why one governed capability, what it covers, and the operating model it follows. This document answers the next question — **who does what, in what order, starting when** — for the compressed two-year timeline the strategy commits to (Year 1: Foundation & Integration; Year 2: Operational & Continuous).

Governance Stand-Up (Weeks 1–4)

Action	Owner	Exit criterion
Charter the AI Assurance Governance Board; name a chair	Sponsoring executive	Charter signed; board seated
Ratify the Risk-Tiered Depth model (Tiers 1–4) as the enterprise standard	Governance Board	Published policy, no open objections from major stakeholders
Inventory every existing evaluation effort (LLM, translation/NER/speech, harness, operational test)	Governance Board + domain leads	Inventory published, mapped to the capability taxonomy
Stand up the shared Enterprise Capability Taxonomy (v1)	Metrics/Standards lead	Taxonomy document ratified by Governance Board

Workstreams (mirrors the Operating Model’s five capabilities)

Workstream	Q1–Q2 (Foundation)	Q3–Q4 (Integration)
Benchmark & Dataset Registries	Stand up shared registries; migrate the highest-value existing benchmarks/datasets first (LLM, translation/NER/speech)	Onboard remaining domain-team benchmarks; retire duplicated per-team copies
Evaluation Orchestration Platform	Select/build submit → select → execute → score → report pipeline; pilot on one domain	Platform live for Model and Harness evaluation enterprise-wide
Model / Harness Evaluation	Stand up AICO (or equivalent) as system of record for model scoring and Bradley-Terry/Glicko-2/GLMM harness comparison	Harness maturity radar and cost/Pareto frontier reporting live; executive scorecards replace per-team reports

Workstream	Q1–Q2 (Foundation)	Q3–Q4 (Integration)
AI-Enabled System Evaluation	Define pipeline-level scoring methodology for the first candidate system (translation → NER → retrieval → LLM → decision)	Roll out to remaining Tier 1–2 systems
Operational & Continuous Monitoring	Design OT&E-style test scenarios and red-team playbooks (MITRE ATLAS-mapped) for Tier 1 systems	Production drift monitoring and scheduled re-certification become the default lifecycle for Tier 1–2

First 90 Days — Concrete Checklist

1. Governance Board chartered and seated (Week 2).
2. Existing evaluation-effort inventory complete, mapped to the capability taxonomy (Week 4).
3. Enterprise Capability Taxonomy v1 ratified (Week 6).
4. One pilot benchmark/dataset registry live, migrated from an existing team’s ad-hoc store (Week 8).
5. First Bradley-Terry harness leaderboard published from real submitted trial data (Week 10).
6. First executive scorecard delivered to sponsoring leadership, replacing a legacy per-team report (Week 12).

Staffing

Role	Count	Sourced from
Governance Board members	5–7	Existing domain leads (LLM, traditional ML, harness, security, operational test) — a standing committee, not new headcount
Benchmark/Dataset Engineering	2–3 FTE	Reassigned from consolidated per-team benchmark efforts
Platform Engineering (orchestration)	2–3 FTE	New or reassigned from existing MLOps/DevOps capacity
Metrics & Standards lead	1 FTE	New — the single owner of the capability taxonomy and metric catalog
Domain evaluators (Model, Harness, AI-Enabled System, Operational)	Retained in place	Existing domain teams keep their expertise; they adopt shared standards, they are not disbanded

Success Metrics (Exit Criteria for Year 1)

- 100% of Tier 1–2 systems evaluated against the shared capability taxonomy, not a bespoke per-team rubric.
- One benchmark/dataset registry, one metric catalog, one report format in active use across every domain evaluated this year.
- Executive scorecards fully replace legacy per-team reports for Tier 1–2 systems.
- Harness comparisons (leaderboard, cost/Pareto frontier, maturity radar) published on a recurring cadence, not a one-off study.

Risks & Mitigations

Risk	Mitigation
Domain teams resist consolidation, perceiving loss of ownership	Governance Board seats domain leads directly; teams keep their technical expertise and day-to-day evaluation work, only the standards/registries/reporting layer centralizes
Orchestration platform build takes longer than Q2	Pilot on the highest-value single domain (Model/Harness evaluation via AICO) first; do not block other workstreams on platform completion
Compressed one-year Foundation+Integration timeline outpaces available staffing	Sequence registry/taxonomy work (cheap, high-leverage) ahead of platform/orchestration build (expensive); a partially-manual pipeline is an acceptable Q1–Q2 state
AI-Enabled System evaluation methodology is genuinely novel (no established public framework, unlike Model/Harness)	Treat as its own follow-on planning effort once Model/Harness evaluation is stable — do not let its uncertainty delay Year 1 exit criteria for the workstreams that already have a clear path

Dependencies

- The **Model** and **Harness** evaluation capabilities have a working first system of record today (AICO), per the Strategy document’s “A Working Instantiation” — Year 1 execution for that workstream is largely a scaling and governance exercise, not new engineering from zero.
- **AI-Enabled System Evaluation** has no equivalent existing implementation and is scoped as a separate, dedicated planning effort once Model/Harness evaluation reaches steady state.