

Enterprise AI Test, Evaluation & Assurance Policy

The ratified Risk-Tiered Depth model and mandatory evaluation lifecycle

August 2026

Enterprise AI Test, Evaluation & Assurance Policy

Status: DRAFT. Pending ratification by the [AI Assurance Governance Board](#), per the [Implementation Plan](#)'s Governance Stand-Up milestone.

Purpose & Scope

This policy makes the [Strategy](#)'s Risk-Tiered Depth model mandatory. It applies to every AI system in scope of the [AI Assurance Governance Board Charter](#): LLMs, traditional ML, AI harnesses, and AI-enabled systems. A system is in scope from the point it is proposed for enterprise, departmental, or mission use – not only at final deployment.

Policy Statement

No system may be deployed at a given Risk Tier without first satisfying that tier's mandatory gates. Depth of required evaluation scales with deployment risk, not with how convenient a lighter-weight evaluation would be for the owning team.

Risk Tiers and Mandatory Gates

Tier	Scope	Mandatory before deployment
1	Mission-critical / customer-facing	Operational test, red team (MITRE ATLAS-mapped), continuous monitoring enabled, re-certification scheduled – plus everything in Tier 2
2	Enterprise-facing system	AI-enabled system evaluation (pipeline/mission-level, not component-only) and harness evaluation where applicable – plus everything in Tier 3
3	Departmental tool	Independent dataset validation (not solely the developing team's own test set) – plus everything in Tier 4
4	Research / experimental	Benchmark suite evaluation against the Enterprise Capability Taxonomy (baseline requirement for every tier)

Tier assignment is proposed by the owning team and confirmed by the Governance Board at intake; a system is re-tiered if its deployment scope changes materially (e.g., a departmental tool promoted to enterprise-facing).

Recurring Requirements (Tier 1-2)

Passing gates once is not sufficient for Tier 1-2 systems. The following are recurring, not one-time:

- **Continuous monitoring:** production drift detection against the published evaluation baseline, per the cadence set by the Governance Board for that system class.
- **Re-certification:** a full re-evaluation on a fixed schedule or upon a material change (model swap, harness version bump, retraining), whichever comes first.

Roles & Responsibilities

Responsibility	Owner
Propose Tier assignment	System-owning team
Confirm Tier assignment	Governance Board
Execute evaluation against the tier's gates	Relevant domain team(s) (Model/Harness, AI-Enabled System, Operational)
Publish results to the shared registry/reporting layer	Domain team executing the evaluation
Review evidence and issue a readiness recommendation	Governance Board
Final deployment decision	Accountable business/mission owner (the Board's recommendation is evidence, not a veto – see the Charter)

Exceptions

A gate may be waived only via the Chair's emergency waiver process defined in the [Governance Board Charter](#), and only for a time-boxed period with a documented remediation plan. A waiver is recorded in the governance audit trail; it does not exempt the system from the gate permanently.

Standards Referenced

This policy operationalizes, and does not duplicate, the frameworks catalogued in the Strategy's "Grounding in Established Practice": NIST AI RMF's MAP function and DoDI 5000.89's tailoring principle are the basis for tiering; MITRE ATLAS is the basis for the Tier 1 red-team requirement; ISO/IEC 42001's PDCA cycle is the basis for the recurring re-certification requirement.

Enforcement

A system deployed without satisfying its tier's mandatory gates is out of policy. The Governance Board escalates an out-of-policy deployment to the sponsoring executive; this policy does not grant the Board independent enforcement authority beyond escalation, consistent with its charter as an advisory and standards body, not an approval gate unto itself.

A **PDF export** of this page is regenerated via `scripts/generate_docs_pdfs.py` whenever this page changes.