AICO

AI-Enabled System Evaluation Report

Configuration quality, cost, and regression status

demo-rocky / demo-baseline

Overall quality 0.44 · 12 trials evaluated

REPORT GENERATED

26 August 2026, 01:00 UTC

EVALUATION RUN

demo-ai-system-run-2

INTERACTION PATTERN

multi-agent

METHODOLOGY

LLM-as-judge scoring, bootstrap confidence intervals

PREPARED FOR

Executive / deployment review

GENERATED FROM DEMO DATA — FOR DESIGN REVIEW ONLY

Executive Summary

The evaluation of the demo-baseline configuration resulted in an overall quality score of 0.44. This setup operates with a mean cost of $0.0091 per trial and an average response time of 2.2 seconds. The system failed the regression check against its published baseline due to significant decreases in groundedness, helpfulness, and the overall score. Because the performance has declined beyond the acceptable tolerance, this configuration should not be deployed in its current state.

Generated by gemma4:26b via the configured Rocky LLM backend.

OVERALL QUALITY	TRIALS EVALUATED	MEAN COST / TRIAL	REGRESSION CHECK
0.44	12	$0.009	Failed

This report profiles demo-baseline using 12 LLM-judged evaluation trials from run demo-ai-system-run-2, comparing it against every other configuration evaluated in the same run and against its own published baseline.

Scores reflect the most recently submitted evaluation run. They will shift as new trials are submitted — treat this snapshot as current-state, not a permanent record.

Metric Breakdown

demo-baseline's mean score on every metric this evaluation scores, with a 95% bootstrap confidence interval across 12 trials.

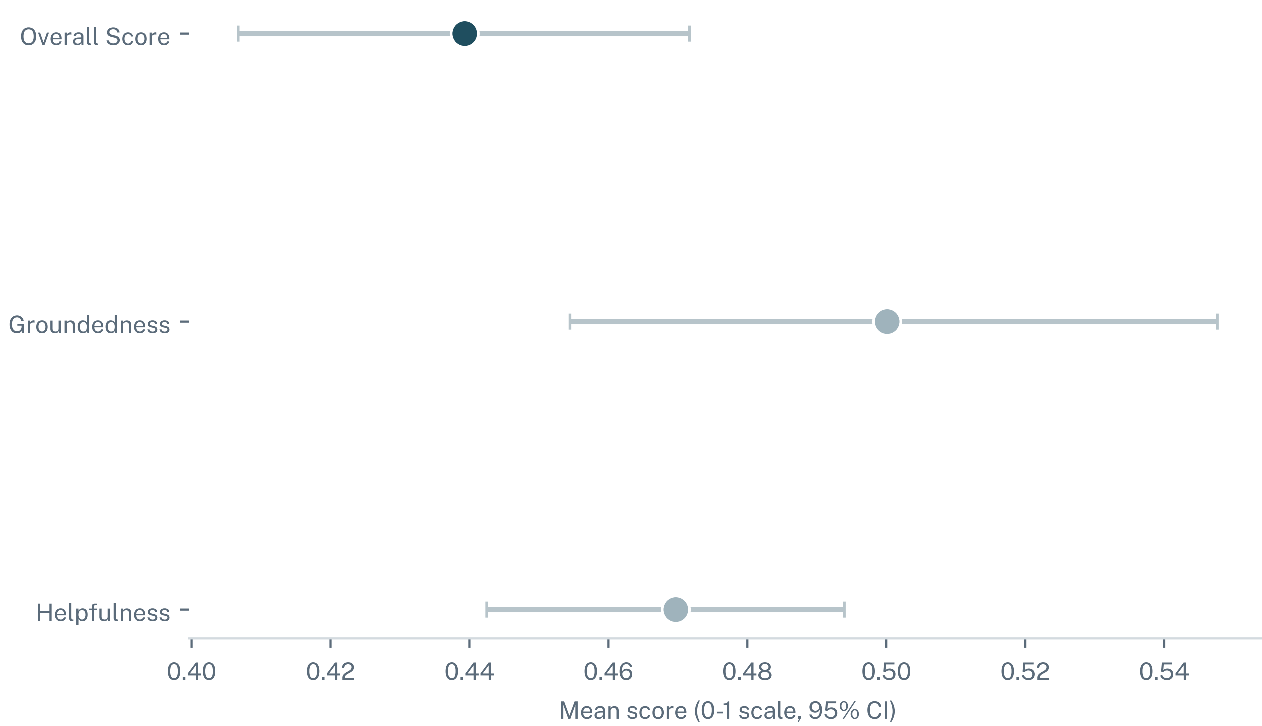

Metric	Mean	95% CI	N trials
Overall Score	0.439	[0.407, 0.472]	12
Groundedness	0.500	[0.454, 0.548]	12
Helpfulness	0.470	[0.442, 0.494]	12

Cost & Efficiency

What it costs to run demo-baseline in production, per trial, based on this evaluation run's real token usage and elapsed time.

MEAN COST / TRIAL	MEAN ELAPSED TIME	MEAN TOKENS IN	MEAN TOKENS OUT
$0.0091	2.2s	609	120

WHAT THIS MEANS

Cost and latency are measured directly from this configuration's own evaluation trials, not estimated — the same token accounting used to bill a production request. A configuration with a higher quality score is only a better choice once its cost per trial is weighed against how much that quality gain is actually worth at deployment volume.

Comparison vs. Other Configurations

Overall quality score for every configuration evaluated in run demo-ai-system-run-2, with demo-baseline highlighted.

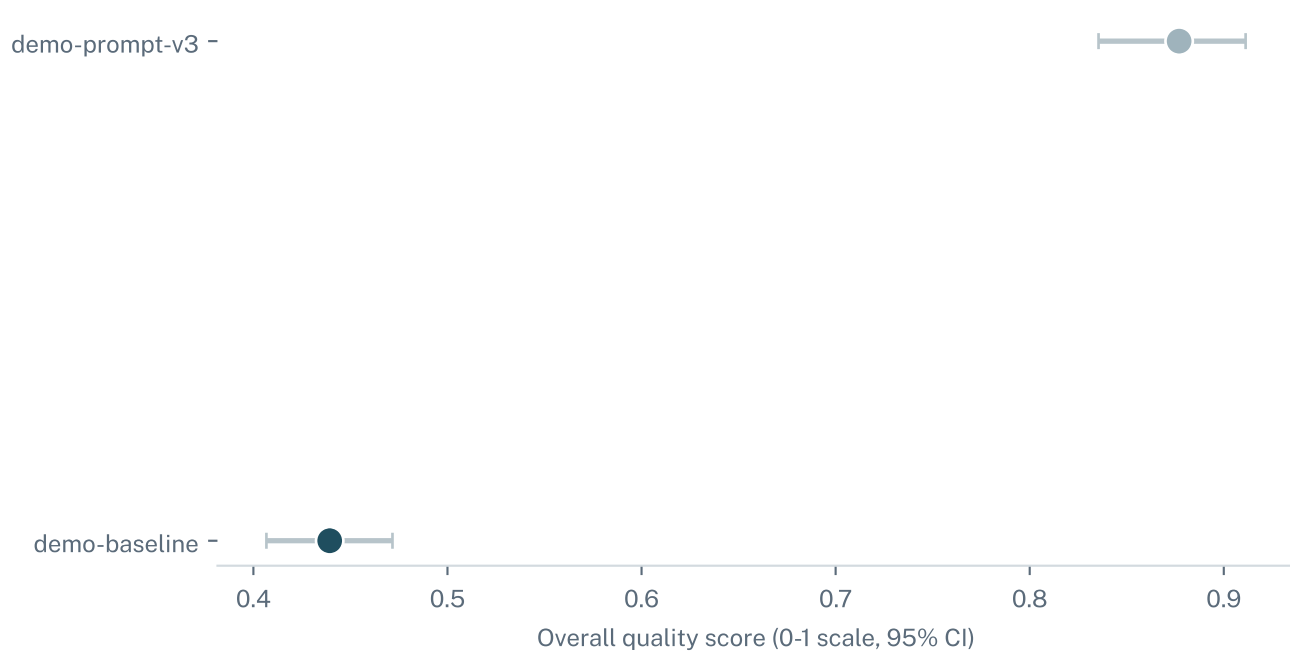

Configuration	Overall quality	95% CI	N trials
demo-prompt-v3	0.877	[0.835, 0.911]	12
demo-baseline	0.439	[0.407, 0.472]	12

Regression vs. Published Baseline

WHAT THIS MEANS

A regression check compares this run's mean scores against a baseline the team published earlier from a trusted run of the same configuration. A metric is 'stable' if it held steady or improved, 'watch' if it slipped a little, and 'regressing' if it dropped by more than the configured tolerance — a governance gate for catching a quality drop (e.g. from a prompt or model change) before it reaches production.

FAILED

12 trials vs. the baseline published from run demo-ai-system-run-1 · tolerance 0.10

Metric	Baseline	Current	Delta	Status
Groundedness	0.706	0.500	-0.206	REGRESSING
Helpfulness	0.637	0.470	-0.168	REGRESSING
Overall Score	0.686	0.439	-0.247	REGRESSING

Methodology & Glossary

LLM-as-judge scoring

Each evaluated trial's transcript is graded by a separate, independent language model against a fixed rubric (helpfulness, groundedness, and an overall quality score) rather than a hand-written rule set, since a conversational response's quality is not easily captured by exact-match or keyword checks.

95% confidence interval

The range this report is 95% confident contains the configuration's true mean score, estimated via bootstrap resampling of the underlying trial scores. A narrow interval means high confidence; a wide one means more trials are needed.

Regression gate

A pass/fail check comparing a new evaluation run's mean scores against a previously published baseline for the same configuration, catching a quality drop before it reaches production. See page 6 for this configuration's own result.

Cost per trial

Computed directly from this configuration's own token usage and elapsed time on each evaluated trial — the same accounting used to bill a real production request, not an estimate.

DATA & METHOD

This report was generated from synthetic demo data for design review, not a real AI-enabled system evaluation submission. Data shown here should not be used for a real deployment decision.