

AICO

Benchmark Evaluation Report

Score distribution, ranking contribution, and governance coverage

sim-bench-0000

Domain sim-domain-000 · mean score 124.8 · 16 models scored

REPORT GENERATED

26 August 2026, 01:00 UTC

CLASSIFICATION

Tier 3 · internal · 2PL IRT model

EVALUATION SCOPE

16 models · 20 benchmarks · 6 domains

PREPARED FOR

Executive / benchmark-portfolio review

GENERATED FROM DEMO DATA — FOR DESIGN REVIEW ONLY

Executive Summary

This report evaluates the sim-bench-0000 benchmark, which measures performance within the sim-domain-000 tier 3 category using internal data. The evaluation included 16 of 16 models, showing a near-perfect correlation of 0.99 with the Global ICI ranking. Because the results track the overall model rankings so closely, this benchmark provides highly reliable performance insights. Given this strong alignment with global standards, it is a valuable metric to continue running for our ongoing assessments.

Generated by gemma4:26b via the configured Rocky LLM backend.

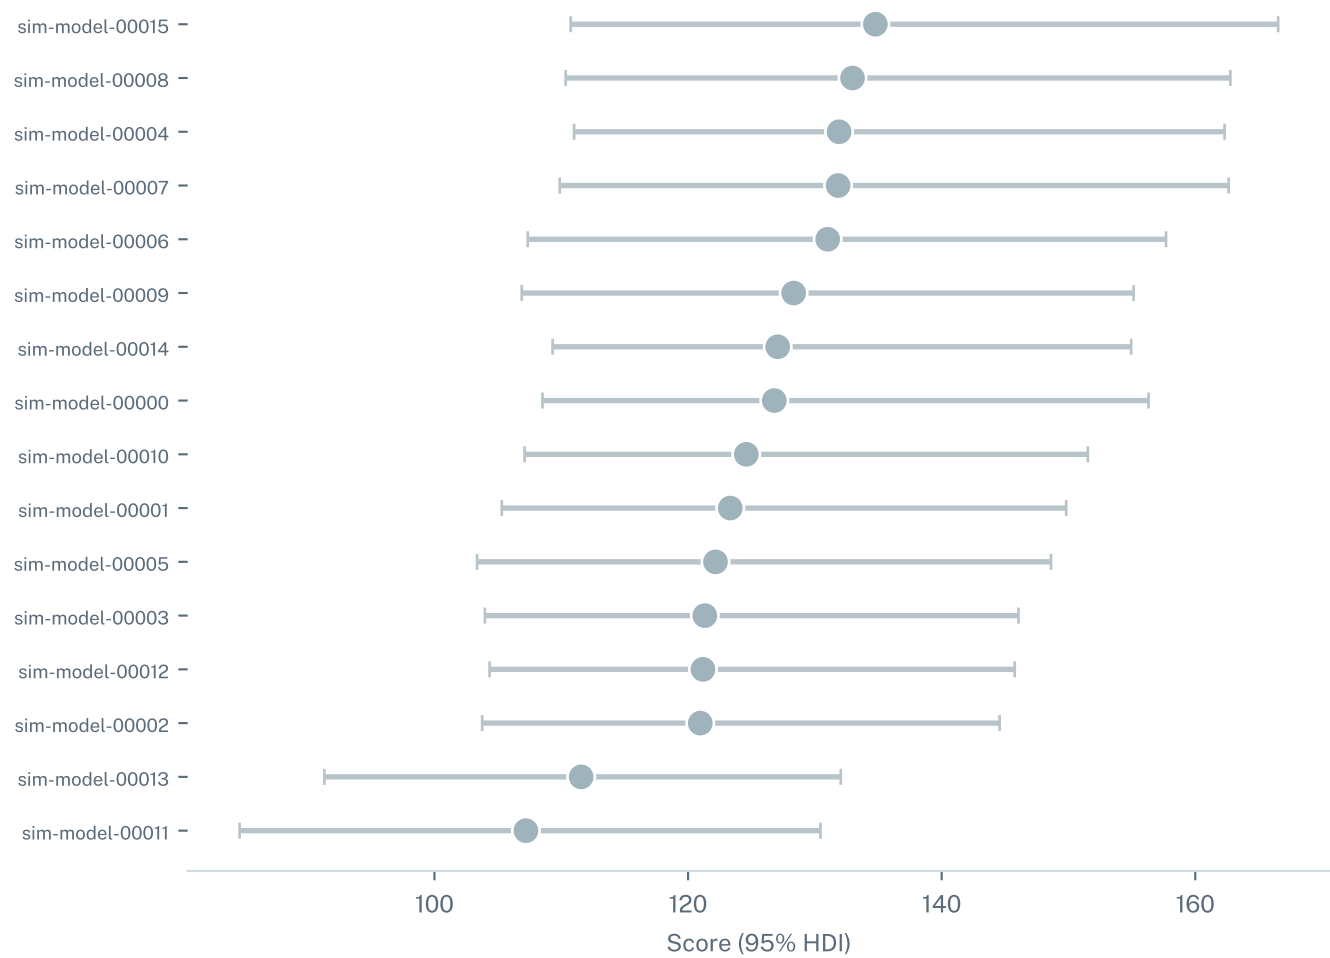
MEAN SCORE	MODELS SCORED	GLOBAL RANK CORRELATION	TIER
124.8	16	0.99	Tier 3

This report profiles sim-bench-0000, one of 20 benchmarks in the evaluated suite, on its own score distribution and how well it tracks each model's overall standing.

Scores reflect the most recently published fit. They will shift as new observations are submitted — treat this snapshot as current-state, not a permanent record.

Model Scores on This Benchmark

Every model's score on sim-bench-0000, with a 95% highest-density credible interval, ranked highest to lowest.



Top model	Mean	95% HDI
sim-model-00015	134.77	[110.75, 166.54]
sim-model-00008	132.97	[110.35, 162.77]
sim-model-00004	131.92	[111.02, 162.31]
sim-model-00007	131.84	[109.88, 162.63]
sim-model-00006	131.01	[107.36, 157.69]
sim-model-00009	128.33	[106.89, 155.14]

Contribution to the Global Ranking

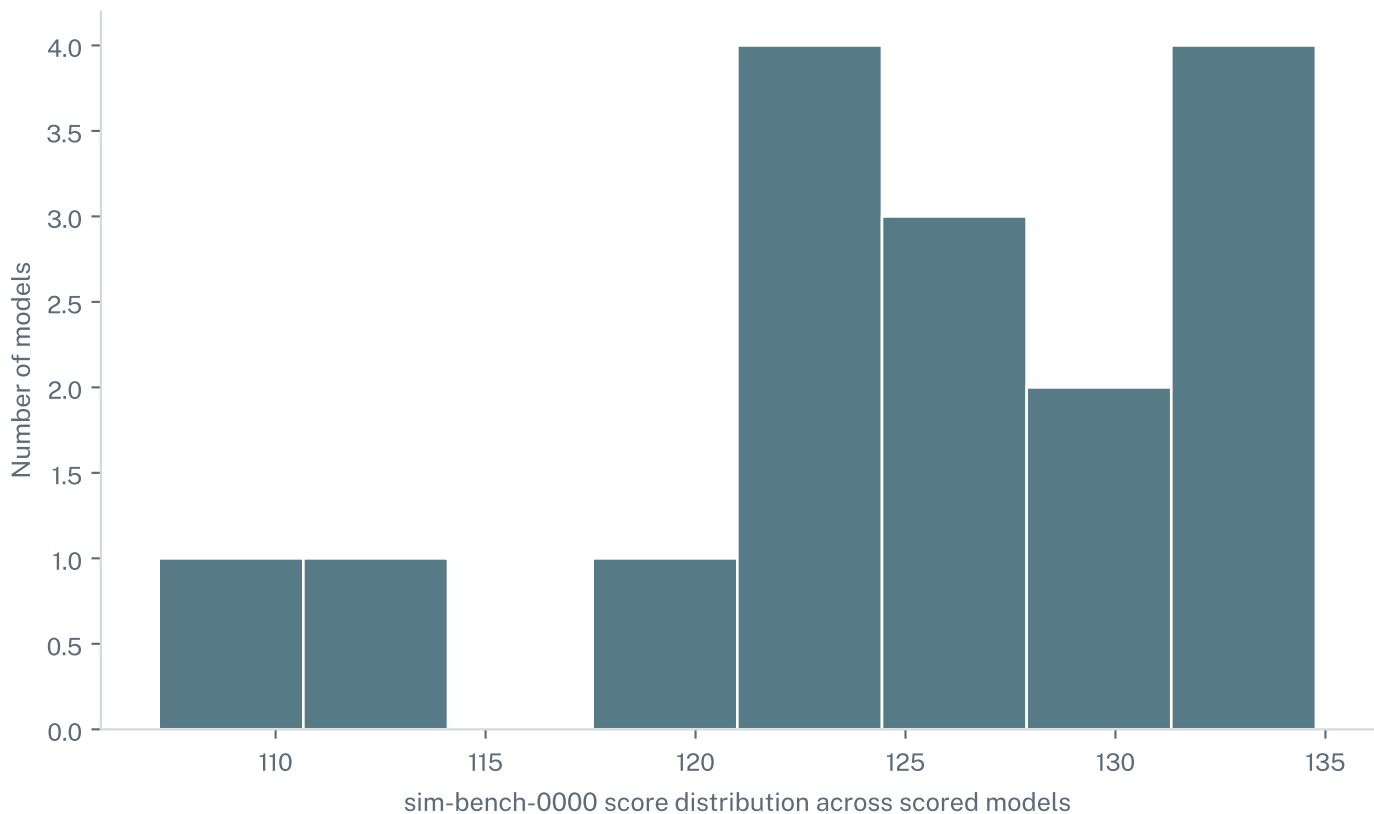
WHAT THIS MEANS

This measures how closely a benchmark's own ranking of models matches the overall Global ICI ranking (a Spearman rank correlation). A benchmark near +1 moves in lockstep with the overall ranking; one near 0 tests something the overall ranking doesn't capture; a negative value would mean it actively disagrees. A benchmark that never diverges from the overall ranking may be worth reviewing for redundancy.

sim-bench-0000's correlation with the Global ICI ranking

0.99

mean score 124.82 · score std. dev 7.33 · 16 models shared with the global ranking



Coverage & Classification

How sim-bench-0000 is classified and tagged for governance and threat-coverage tracking.

CLASSIFICATION

Domain: sim-domain-000
Tier: Tier 3
Source: internal
IRT model: 2PL

OWNERSHIP

Organization: Acme Corp
System: Assistant Platform

MITRE ATLAS TECHNIQUES COVERED

- AML.T0000 — Search Open Technical Databases
- AML.T0000.000 — Journals and Conference Proceedings
- AML.T0000.001 — Pre-Print Repositories

Portfolio Redundancy Signal

WHAT THIS MEANS

Two benchmarks are redundant when they tend to rank models the same way — a high normalized mutual information (NMI) score means one benchmark's result rarely tells you anything the other didn't already. A portfolio with several highly redundant benchmarks spends evaluation budget re-measuring the same signal instead of covering new ground; this doesn't mean a benchmark should be retired on its own, only that it is worth a second look alongside its most-similar peers.

Most similar benchmark	Normalized mutual information
sim-bench-0016	0.37
sim-bench-0015	0.36
sim-bench-0002	0.35

Methodology & Glossary

95% highest-density interval (HDI)

The range this report is 95% confident contains a model's true score on this benchmark, given the data evaluated so far. A narrow interval means high confidence; a wide one means the score should be treated as preliminary.

Spearman rank correlation

How closely this benchmark's own ranking of models matches the overall Global ICI ranking. See page 4 for this benchmark's own result.

Normalized mutual information (NMI)

A measure of how redundant two benchmarks are with each other, on a 0-1 scale. See page 6 for this benchmark's most similar peers.

MITRE ATLAS technique

A named adversarial ML technique from the MITRE ATLAS knowledge base that this benchmark has been manually tagged as covering, for threat-coverage tracking.

DATA & METHOD

This report was generated from a synthetic demo ecosystem for design review, not a real benchmark evaluation. Organization/system and MITRE tags shown are illustrative and should not be used for a real benchmark-portfolio decision.