

Enterprise Capability Taxonomy

The common vocabulary every benchmark, dataset, score, and report maps to

August 2026

Enterprise Capability Taxonomy

Status: DRAFT (v1). Pending ratification by the [AI Assurance Governance Board](#), per the [Implementation Plan](#)’s Governance Stand-Up milestone.

Purpose

Every benchmark, dataset, score, and report produced under the [Enterprise AI Test, Evaluation & Assurance Strategy](#) maps to exactly one entry in this taxonomy. Without a shared vocabulary, “reasoning” means something different to each team that measures it, and scores stop being comparable the moment the underlying models change. This taxonomy is the thing that stays stable while everything under it churns.

Structure

Two levels: a small set of top-level **capability domains**, each with **sub-capabilities** that a benchmark or dataset can target directly. A benchmark maps to exactly one sub-capability; a system-level (AI-Enabled System or Operational) evaluation may span several.

Domain	Sub-capabilities	Example metric family
Reasoning	Multi-step problem solving, mathematical reasoning, logical inference	Accuracy, chain-of-thought correctness
Coding	Code generation, code editing, debugging, code review	Pass@k, test-suite pass rate
Translation	Machine translation, cross-lingual retrieval	BLEU, COMET, chrF
NER	Named-entity recognition, entity linking	Precision/recall/F1 (CoNLL-style)
Speech	Speech-to-text, speaker identification	WER (word error rate)
Vision	Image classification, object detection, visual QA	Accuracy, mAP
Forecasting	Time-series prediction, trend estimation	MAE, coverage calibration
Harness Quality	Tool use, task completion, cost/time efficiency, operational profile	Bradley-Terry rating, pass rate, cost/token efficiency, five-axis maturity profile (see Data Inputs, Outputs & Measurement Points)
System/Mission Outcomes	End-to-end pipeline effectiveness, human-AI teaming, mission task completion	Mission-specific KPI, time-to-decision, error rate reduction

Governance of the Taxonomy

- **Additions:** a new sub-capability may be proposed by any domain team; the Metrics & Standards lead reviews for overlap with an existing entry before it reaches the Governance Board for ratification.
- **No silent proliferation:** a proposed addition that substantially overlaps an existing sub-capability is folded into it rather than added as a new one – the taxonomy’s value is in staying small enough to be a genuine shared vocabulary, not an ever-growing list that re-fragments into per-team dialects.
- **Versioning:** this document is versioned (v1 above); a breaking change to an existing domain’s definition requires Governance Board ratification and a version bump, the same discipline `aico.harness.contracts`’ `schema_version` field applies to its own submission contract.

Mapping Convention

Every benchmark registered anywhere in the evaluated estate declares which sub-capability it targets. In AICO specifically, this is the existing `domain_id` field already carried on every `BenchmarkMetadata` record (POST `/benchmarks`) – adopting this taxonomy means populating that field from the ratified sub-capability list above, rather than each team choosing its own domain names. AI harness evaluation maps to **Harness Quality** regardless of which underlying model drove the trial, consistent with the Strategy’s “separate harness from model” objective: a harness’s Bradley-Terry rating, cost/Pareto position, and maturity profile are Harness Quality measurements, never folded into the model’s own Reasoning/Coding scores.

Relationship to Risk Tiers

The taxonomy answers “what is being measured”; the [T&E Policy](#)’s Risk Tiers answer “how much evidence is required before deployment.” The two are independent: a Tier 1 system’s Coding capability and a Tier 4 research prototype’s Coding capability are measured against the *same* sub-capability definition and metric family – only the required evidence depth differs.

A [PDF export](#) of this page is regenerated via `scripts/generate_docs_pdfs.py` whenever this page changes.