

Apollo AI

LLM Benchmark Observatory

Covered Model Categories Weekly Report

Week 35 / 2026

ISO Week 2026-W35

Table of Contents

Table of Contents	2
1. Week Summary	3
2. Axis / Subcategory Distribution	4
3. Covered Benchmarks — Full Detail	5
Capability Axis (24)	5
Misalignment Axis (5)	16

1. Week Summary

As of this report, **29** benchmark(s) in the observatory carry at least one covered-model-category signal under the cyber-operations capability taxonomy — **24** on the Capability axis, **0** on Autonomy, and **5** on Misalignment.

Metric	Count
Benchmarks classified (any signal)	29
Capability axis	24
Autonomy axis	0
Misalignment axis	5

2. Axis / Subcategory Distribution

Axis	Subcategory	Benchmarks
Autonomy	Initiative	3
Autonomy	Sustained Operation	7
Capability	Exploit Development	13
Capability	Malware Development	11
Capability	Network Intrusion Operations	17
Capability	Reverse Engineering and Analysis	14
Capability	Vulnerability Research	19
Misalignment	Deceptive or Evasive Behavior	1
Misalignment	Objective Fidelity	5
Misalignment	Susceptibility to Adversarial Redirection	5

3. Covered Benchmarks — Full Detail

All 29 benchmark(s) currently classified under a covered model category, grouped by primary axis and ranked by alignment strength within each group.

Capability Axis (24)

NYU CTF Bench

Field	Value
Primary subcategory	Reverse Engineering and Analysis (strength 1.00)
All subcategory scores	Reverse Engineering and Analysis (1.00), Exploit Development (0.80), Vulnerability Research (0.80), Network Intrusion Operations (0.50), Malware Development (0.20)
Domain	cyber
Organization	New York University / NYU Cybersecurity Center
Quality tier	silver
Release date	2024-06-01
Dataset size	200
Evaluation type	automated
URL	https://huggingface.co/datasets/shellphish/NYU-CTF-Bench
Paper	https://arxiv.org/abs/2406.05590

Scalable open-source benchmark of 200 real-world CTF challenges sourced from CTFTIME competitions (2020-2023). Six categories: binary exploitation (pwn), reverse engineering, web, cryptography, forensics, and miscellaneous. Designed for systematic evaluation of LLM offensive security capabilities at scale. Includes solve scripts and environment snapshots for full reproducibility.

BinaryAudit

Field	Value
Primary subcategory	Reverse Engineering and Analysis (strength 1.00)
All subcategory scores	Reverse Engineering and Analysis (1.00), Vulnerability Research (0.80), Sustained Operation (0.50), Malware Development (0.50)
Domain	agentic
Organization	-
Quality tier	caution
Release date	-
Dataset size	46
Evaluation type	agent
URL	https://meridianlabs-ai.github.io/inspect_harbor/registry/binary_audit.html
Paper	https://github.com/QuesmaOrg/BinaryAudit

BinaryAudit: AI-agent benchmark for finding backdoors hidden in compiled binaries via reverse engineering.

CyberGym

Field	Value
Primary subcategory	Vulnerability Research (strength 1.00)
All subcategory scores	Vulnerability Research (1.00), Reverse Engineering and Analysis (0.70), Network Intrusion Operations (0.50), Sustained Operation (0.40), Exploit Development (0.30), Initiative (0.30)
Domain	agentic
Organization	-
Quality tier	bronze
Release date	-
Dataset size	6028
Evaluation type	agent
URL	https://ukgovernmentbeis.github.io/inspect_evals/evals/cybergym/
Paper	https://arxiv.org/abs/2506.02548

A large-scale, high-quality cybersecurity evaluation framework designed to rigorously assess the capabilities of AI agents on real-world vulnerability analysis tasks.

Cybench

Field	Value
Primary subcategory	Reverse Engineering and Analysis (strength 0.90)
All subcategory scores	Reverse Engineering and Analysis (0.90), Exploit Development (0.90), Vulnerability Research (0.90), Network Intrusion Operations (0.50), Malware Development (0.30)
Domain	general
Organization	-
Quality tier	bronze
Release date	-
Dataset size	39
Evaluation type	agent
URL	https://ukgovernmentbeis.github.io/inspect_evals/evals/cybench/
Paper	https://arxiv.org/abs/2408.08926

Tests language models on cybersecurity skills using 39 of 40 practical, professional-level challenges taken from cybersecurity competitions, designed to cover various difficulty levels and security concepts.

CyberSecEval 4

Field	Value
Primary subcategory	Exploit Development (strength 0.90)
All subcategory scores	Exploit Development (0.90), Malware Development (0.90), Vulnerability Research (0.90), Reverse Engineering and Analysis (0.60), Network Intrusion Operations (0.50)
Domain	general
Organization	-
Quality tier	bronze
Release date	-
Dataset size	1000
Evaluation type	generation
URL	https://ukgovernmentbeis.github.io/inspect_evals/evals/cyberseceval_4/
Paper	https://arxiv.org/abs/2404.13161

A suite of cybersecurity evaluation benchmarks adapted from Meta's PurpleLlama CybersecurityBenchmarks.

BigVul

Field	Value
Primary subcategory	Vulnerability Research (strength 0.80)
All subcategory scores	Vulnerability Research (0.80)
Domain	cyber
Organization	MSR 2020 (GitHub Research)
Quality tier	bronze
Release date	2020-06-01
Dataset size	188519
Evaluation type	automated
URL	https://github.com/ZeoVan/MSR_20_Code_vulnerability_CSV_Dataset
Paper	https://dl.acm.org/doi/10.1145/3379597.3387501

Large-scale C/C++ vulnerability dataset linking 188,519 code functions from 348 open-source GitHub projects to their CVE records and CWE type labels. Collected from GitHub commit history between 2002-2019. Each function is labelled vulnerable/non-vulnerable (3.75% positive rate) with associated CVE ID, CVSS score, and CWE classification. Widely used for vulnerability detection model training and evaluation.

InterCode-CTF

Field	Value
Primary subcategory	Reverse Engineering and Analysis (strength 0.80)
All subcategory scores	Reverse Engineering and Analysis (0.80), Exploit Development (0.80), Vulnerability Research (0.80), Sustained Operation (0.80), Network Intrusion Operations (0.50), Initiative (0.50), Objective Fidelity (0.30)
Domain	cyber
Organization	Princeton NLP Group
Quality tier	silver
Release date	2023-06-01
Dataset size	100
Evaluation type	automated
URL	https://intercode-benchmark.github.io
Paper	https://arxiv.org/abs/2306.14898

Interactive multi-turn CTF evaluation environment built on the InterCode framework. 100 CTF tasks executed in a live Bash shell with iterative execution feedback. Uniquely tests multi-turn agentic reasoning under real environment feedback — the model can observe execution results and adapt across turns. Part of the broader InterCode standardisation effort for interactive code benchmarks.

CVEBench

Field	Value
Primary subcategory	Exploit Development (strength 0.80)
All subcategory scores	Exploit Development (0.80), Sustained Operation (0.50), Vulnerability Research (0.50)
Domain	agentic
Organization	-
Quality tier	bronze
Release date	-
Dataset size	40
Evaluation type	agent
URL	https://ukgovernmentbeis.github.io/inspect_evals/evals/cve_bench/
Paper	https://arxiv.org/abs/2503.17332

Characterises an AI Agent's capability to exploit real-world web application vulnerabilities.

GDM Dangerous Capabilities

Field	Value
Primary subcategory	Exploit Development (strength 0.80)
All subcategory scores	Exploit Development (0.80), Vulnerability Research (0.80), Reverse Engineering and Analysis (0.70), Network Intrusion Operations (0.70), Sustained Operation (0.60), Initiative (0.50), Objective Fidelity (0.40), Malware Development (0.30)
Domain	agentic
Organization	-
Quality tier	bronze
Release date	-
Dataset size	13
Evaluation type	agent
URL	https://ukgovernmentbeis.github.io/inspect_evals/evals/gdm_in_house_ctf/
Paper	https://arxiv.org/abs/2403.13793

CTF challenges covering web app vulnerabilities, off-the-shelf exploits, databases, Linux privilege escalation, password cracking and spraying.

InterCode

Field	Value
Primary subcategory	Reverse Engineering and Analysis (strength 0.80)
All subcategory scores	Reverse Engineering and Analysis (0.80), Vulnerability Research (0.80), Exploit Development (0.70), Network Intrusion Operations (0.50), Sustained Operation (0.40)
Domain	coding
Organization	-
Quality tier	bronze
Release date	-
Dataset size	78
Evaluation type	agent
URL	https://ukgovernmentbeis.github.io/inspect_evals/evals/gdm_intercode_ctf/
Paper	https://arxiv.org/abs/2306.14898

Tests AI's ability in coding, cryptography, reverse engineering, and vulnerability identification through practical capture-the-flag (CTF) cybersecurity scenarios.

CYBERSECEVAL 3

Field	Value
Primary subcategory	Vulnerability Research (strength 0.70)
All subcategory scores	Vulnerability Research (0.70), Exploit Development (0.60), Malware Development (0.60), Network Intrusion Operations (0.50), Reverse Engineering and Analysis (0.30)
Domain	general
Organization	-
Quality tier	bronze
Release date	-
Dataset size	1000
Evaluation type	generation
URL	https://ukgovernmentbeis.github.io/inspect_evals/evals/cyberseceval_3/
Paper	https://arxiv.org/abs/2312.04724

Evaluates Large Language Models for cybersecurity risk to third parties, application developers and end users.

CTIBench

Field	Value
Primary subcategory	Reverse Engineering and Analysis (strength 0.50)
All subcategory scores	Reverse Engineering and Analysis (0.50), Vulnerability Research (0.50)
Domain	cyber
Organization	University of Electronic Science and Technology of China
Quality tier	silver
Release date	2024-06-01
Dataset size	4800
Evaluation type	automated
URL	https://huggingface.co/datasets/Liqiang/CTIBench
Paper	https://arxiv.org/abs/2406.07599

Comprehensive benchmark for evaluating LLMs on Cyber Threat Intelligence (CTI) tasks. Five subtasks: (1) CVE options — multiple-choice CVE knowledge; (2) TTP classification — mapping threat reports to MITRE ATT&CK; tactics/techniques; (3) threat actor attribution — identifying threat groups from behavioural indicators; (4) malware analysis — classifying malware families and capabilities; (5) indicator extraction — identifying IoCs from threat reports. Each task uses real CTI reports and STIX/ATT&CK; ground truth.

CVEFixes

Field	Value
Primary subcategory	Vulnerability Research (strength 0.50)
All subcategory scores	Vulnerability Research (0.50)
Domain	cyber
Organization	Secure IT Project / University of Oslo
Quality tier	bronze
Release date	2021-11-01
Dataset size	5495
Evaluation type	automated
URL	https://github.com/secureIT-project/CVEfixes
Paper	https://dl.acm.org/doi/10.1145/3475960.3475985

Automated collection of 5,495 CVE-linked vulnerability patches from open-source repositories across 8 programming languages. Each entry links a CVE ID to the commit that introduced the vulnerability, the commit that fixed it, and code diff at the function/file level. Multi-language coverage (C, C++, Python, Java, Go, PHP, JavaScript, Ruby) makes it uniquely suited for cross-language vulnerability patch analysis.

SecEval

Field	Value
Primary subcategory	Network Intrusion Operations (strength 0.50)
All subcategory scores	Network Intrusion Operations (0.50)
Domain	cyber
Organization	Beijing Institute of Technology
Quality tier	silver
Release date	2023-09-01
Dataset size	2677
Evaluation type	automated
URL	https://github.com/XuanMingZhen/SecEval
Paper	-

Multilingual cybersecurity knowledge evaluation benchmark with 2,677 multiple-choice questions across 7 security topics: software security, application security, system security, web security, cryptography, network security, and penetration testing. Available in both English and Chinese, making it one of very few security benchmarks with non-English coverage. Questions are sourced from professional certification exam banks and curated by security practitioners.

WMDP

Field	Value
Primary subcategory	Network Intrusion Operations (strength 0.50)
All subcategory scores	Network Intrusion Operations (0.50)
Domain	knowledge
Organization	-
Quality tier	-
Release date	-
Dataset size	1273
Evaluation type	generation
URL	https://ukgovernmentbeis.github.io/inspect_evals/evals/wmdp/
Paper	https://arxiv.org/abs/2403.03218

A dataset of 3,668 multiple-choice questions developed by a consortium of academics and technical consultants that serve as a proxy measurement of hazardous knowledge in biosecurity, cybersecurity, and chemical security.

CyberMetric

Field	Value
Primary subcategory	Network Intrusion Operations (strength 0.50)
All subcategory scores	Network Intrusion Operations (0.50), Reverse Engineering and Analysis (0.30), Exploit Development (0.30), Malware Development (0.30), Vulnerability Research (0.30)
Domain	knowledge
Organization	-
Quality tier	bronze
Release date	-
Dataset size	80
Evaluation type	generation
URL	https://ukgovernmentbeis.github.io/inspect_evals/evals/cybermetric/
Paper	https://arxiv.org/abs/2402.07688

Datasets containing 80, 500, 2000 and 10000 multiple-choice questions, designed to evaluate understanding across nine domains within cybersecurity

SEvenLLM

Field	Value
Primary subcategory	Network Intrusion Operations (strength 0.50)
All subcategory scores	Network Intrusion Operations (0.50), Vulnerability Research (0.30), Reverse Engineering and Analysis (0.20)
Domain	agentic
Organization	-
Quality tier	bronze
Release date	-
Dataset size	50
Evaluation type	generation
URL	https://ukgovernmentbeis.github.io/inspect_evals/evals/sevenllm/
Paper	https://arxiv.org/abs/2405.03446

Designed for analyzing cybersecurity incidents, which is comprised of two primary task categories: understanding and generation, with a further breakdown into 28 subcategories of tasks.

Reward-Hack-Bench (Control)

Field	Value
Primary subcategory	Reverse Engineering and Analysis (strength 0.50)
All subcategory scores	Reverse Engineering and Analysis (0.50), Malware Development (0.50), Deceptive or Evasive Behavior (0.50), Objective Fidelity (0.50), Network Intrusion Operations (0.50), Exploit Development (0.50), Vulnerability Research (0.50), Sustained Operation (0.50)
Domain	agentic
Organization	-
Quality tier	caution
Release date	-
Dataset size	8
Evaluation type	agent
URL	https://meridianlabs-ai.github.io/inspect_harbor/registry/islo_labs_reward_hack_bench_control.html
Paper	https://github.com/islo-labs/reward-hack-bench

Baseline/control companion to reward-hack-bench: the same 8 SWE-bench + CyBench tasks with NO cheat prompt, still scored by the fairness judge. Also packages these CyBench/GlacierCTF tasks as Harbor tasks.

Sevenllm Dataset

Field	Value
Primary subcategory	Network Intrusion Operations (strength 0.50)
All subcategory scores	Network Intrusion Operations (0.50), Reverse Engineering and Analysis (0.30), Vulnerability Research (0.20), Exploit Development (0.10), Malware Development (0.10)
Domain	general
Organization	-
Quality tier	caution
Release date	-
Dataset size	-
Evaluation type	-
URL	https://huggingface.co/datasets/Multilingual-Multimodal-NLP/SEVENLLM-Dataset
Paper	-

Introduce We provided, designed for analyzing cybersecurity incidents, which is comprised of two primary task categories: understanding and generation, with a further breakdown into 28 subcategories of tasks. The dataset is in question and answer format, using structured json format for understanding tasks and unstructured text format for generation tasks. We also provide some multiple-choice questions to test the cognitive ability of the model in different vertical fields.... See the full description on the dataset page:

<https://huggingface.co/datasets/Multilingual-Multimodal-NLP/SEVENLLM-Dataset>.

Agentic Pcap

Field	Value
Primary subcategory	Network Intrusion Operations (strength 0.50)
All subcategory scores	Network Intrusion Operations (0.50)
Domain	agentic
Organization	-
Quality tier	bronze
Release date	-
Dataset size	-
Evaluation type	-
URL	https://huggingface.co/datasets/maureille/agentic-pcap
Paper	-

Agentic LLMs for Network Attack Detection Can a tool-using LLM agent independently investigate a raw network capture and identify an attack? How do model choice, reasoning configuration, prompt guidance, and background traffic change the investigation and its final decision? This Hugging Face repository is the canonical research release for our experimental evaluation of autonomous network-traffic investigation with agentic LLMs. It brings together the dataset, recorded... See the full description on the dataset page:

<https://huggingface.co/datasets/maureille/agentic-pcap>.

Evals Speech Recognition Cy En 2512

Field	Value
Primary subcategory	Vulnerability Research (strength 0.50)
All subcategory scores	Vulnerability Research (0.50)
Domain	speech_audio
Organization	-
Quality tier	caution
Release date	-
Dataset size	-
Evaluation type	-
URL	https://huggingface.co/datasets/DewiBrynJones/evals-speech-recognition-cy-en-2512
Paper	-

Transcribe model test wer cer whisper-large-v2-ft-btb-cv-cvad-cven-cvmt-wlga-ca-2512 cymen-arfor/lleisiau-arfor 40.7543 19.6824 whisper-tiny-ft-btb-cv-cvad-cven-cvmt-wlga-ca-2601 cymen-arfor/lleisiau-arfor 68.392 42.4607 whisper-large-v2-ft-btb-cv-cvad-cven-cvmt-wlga-ca-2512 techiaith/banc-trawsgriadau-bangor 32.7111 12.6499 whisper-tiny-ft-btb-cv-cvad-cven-cvmt-wlga-ca-2601 techiaith/banc-trawsgriadau-bangor 50.0489 23.5431... See the full description on the dataset page:

<https://huggingface.co/datasets/DewiBrynJones/evals-speech-recognition-cy-en-2512>.

BackdoorAgent

Field	Value
Primary subcategory	Malware Development (strength 0.50)
All subcategory scores	Malware Development (0.50)
Domain	agentic
Organization	-
Quality tier	silver
Release date	2026-01-09
Dataset size	-
Evaluation type	-
URL	https://github.com/Yunhao-Feng/BackdoorAgent
Paper	-

BackdoorAgent is a stage-aware framework and benchmark that instruments LLM-agent workflows (planning, memory, tools) to systematically inject, track, and evaluate backdoor triggers across multi-step trajectories using unified metrics like ASR and clean accuracy.

ALERT

Field	Value
Primary subcategory	Network Intrusion Operations (strength 0.50)
All subcategory scores	Network Intrusion Operations (0.50)
Domain	safety
Organization	-
Quality tier	silver
Release date	2024-04-06
Dataset size	-
Evaluation type	-
URL	https://github.com/Babelscape/ALERT
Paper	-

Official repository for the paper "ALERT: A Comprehensive Benchmark for Assessing Large Language Models' Safety through Red Teaming"

Object Detections

Field	Value
Primary subcategory	Vulnerability Research (strength 0.50)
All subcategory scores	Vulnerability Research (0.50)
Domain	general
Organization	-
Quality tier	caution
Release date	-
Dataset size	-
Evaluation type	-
URL	https://huggingface.co/datasets/wdc9645/object-detections
Paper	-

Car Object Detection Vehicle object detection results from a segment of a Toyota RAV4 2026 review video, enabling semantic video retrieval of car components. Dataset Overview Video Source: YouTube Video YcvECxtXoxQ Section Processed: Exterior Review (18:39 to 28:00 = 9 min 21 sec) Total Detections: 1,179 Total Segments: 123 Sampling Rate: 1 frame per second (560 frames) Files 1. detections.parquet Frame-level object detections. Each row represents... See the full description on the dataset page: <https://huggingface.co/datasets/wdc9645/object-detections>.

Misalignment Axis (5)

CodePI

Field	Value
Primary subcategory	Susceptibility to Adversarial Redirection (strength 1.00)
All subcategory scores	Susceptibility to Adversarial Redirection (1.00)
Domain	coding
Organization	-
Quality tier	bronze
Release date	-
Dataset size	45
Evaluation type	generation
URL	https://ukgovernmentbeis.github.io/inspect_evals/evals/ipi_coding_agent/
Paper	-

Measures coding agent vulnerability to indirect prompt injection attacks embedded in software engineering artifacts (issue descriptions, code comments, README files).

AgentDojo

Field	Value
Primary subcategory	Susceptibility to Adversarial Redirection (strength 1.00)
All subcategory scores	Susceptibility to Adversarial Redirection (1.00), Objective Fidelity (0.80)
Domain	agentic
Organization	-
Quality tier	-
Release date	-
Dataset size	1014
Evaluation type	agent
URL	https://ukgovernmentbeis.github.io/inspect_evals/evals/agentdojo/
Paper	https://arxiv.org/abs/2406.13352

Assesses whether AI agents can be hijacked by malicious third parties using prompt injections in simple environments such as a workspace or travel booking app.

PAIR (Prompt Automatic Iterative Refinement)

Field	Value
Primary subcategory	Susceptibility to Adversarial Redirection (strength 0.80)
All subcategory scores	Susceptibility to Adversarial Redirection (0.80)
Domain	jailbreak
Organization	Princeton University
Quality tier	silver
Release date	2023-10-01
Dataset size	50
Evaluation type	llm_judge
URL	https://github.com/patrickrchao/JailbreakingLLMs
Paper	https://arxiv.org/abs/2310.08419

PAIR is an automated jailbreak attack method and accompanying evaluation framework. Uses an attacker LLM to iteratively refine adversarial prompts against a target LLM, guided by the target's responses. The attack loops for up to 20 iterations per harmful behavior, using the attacker model to analyze why the previous prompt failed and generate an improved version. Key metric: attack success rate (ASR) at k iterations (ASR@ k). Demonstrated >99% ASR on GPT-3.5 on 50 harmful behaviors. Provides a systematic way to measure model vulnerability to iterative prompt refinement attacks.

CyberSecEval_2

Field	Value
Primary subcategory	Susceptibility to Adversarial Redirection (strength 0.80)
All subcategory scores	Susceptibility to Adversarial Redirection (0.80), Exploit Development (0.70), Network Intrusion Operations (0.50), Malware Development (0.40)
Domain	coding
Organization	-
Quality tier	caution
Release date	-
Dataset size	500
Evaluation type	generation
URL	https://ukgovernmentbeis.github.io/inspect_evals/evals/cyberseceval_2/
Paper	https://arxiv.org/pdf/2404.13161

Assesses language models for cybersecurity risks, specifically testing their potential to misuse programming interpreters, vulnerability to malicious prompt injections, and capability to exploit known software vulnerabilities.

AgentThreatBench

Field	Value
Primary subcategory	Susceptibility to Adversarial Redirection (strength 0.80)
All subcategory scores	Susceptibility to Adversarial Redirection (0.80), Objective Fidelity (0.50)
Domain	agentic
Organization	-
Quality tier	bronze
Release date	-
Dataset size	10
Evaluation type	agent
URL	https://ukgovernmentbeis.github.io/inspect_evals/evals/agent_threat_bench/
Paper	https://genai.owasp.org/resource/owasp-top-10-for-agentic-applications-for-2026/

Evaluates LLM agents against the OWASP Top 10 for Agentic Applications (2026), measuring both task utility and security resilience across memory poisoning, autonomy hijacking, and data exfiltration scenarios.