AICO

AI Harness Evaluation Report

Coding-agent harness performance & competitive positioning

demo-claude-code

Bradley-Terry rank 1 of 3 · rating 1.25

REPORT GENERATED

26 August 2026, 01:00 UTC

EVALUATION SCOPE

3 harnesses · 3 models · 5 tasks · 135 trials

METHODOLOGY

Bradley-Terry pairwise comparison, nonparametric bootstrap

PREPARED FOR

Executive / tooling procurement review

GENERATED FROM DEMO DATA — FOR DESIGN REVIEW ONLY

Executive Summary

The demo-claude-code coding harness ranked first among the three systems evaluated in this study. This harness achieved an 87 percent success rate at an average cost of 0.237 dollars per trial, representing an ideal balance where no other system provided higher accuracy for a lower price. Because this system provides the most efficient trade-off between performance and expenditure, it is the recommended choice for deploying automated development tasks.

Generated by gemma4:26b via the configured Rocky LLM backend.

BT RANK	BT RATING	MEAN PASS RATE	MEAN COST / TRIAL
1 / 3	1.25	87%	$0.237

This report profiles demo-claude-code against every other harness evaluated on the same (model, task) combinations — a Bradley-Terry comparison, the same method behind chess and chatbot-arena style rankings, chosen because an agentic coding trial is a multi-step outcome, not a single test-taker response an IRT model could score directly.

Rating and rank reflect the most recently submitted trial results. Both will shift as new trials, models, or harnesses are added — treat this snapshot as current-state, not a permanent standing.

Competitive Positioning

Bradley-Terry rating vs. every other harness demo-claude-code was compared against, with a 95% bootstrap confidence interval.

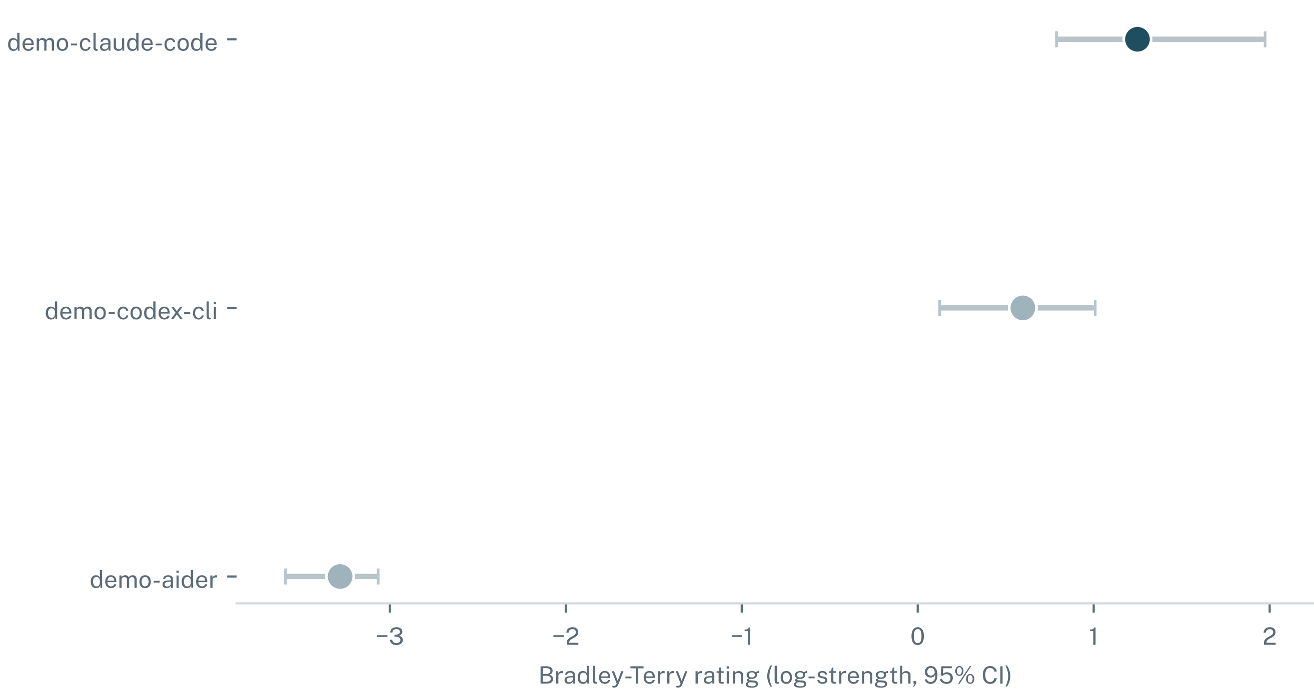

Rank	Harness	Rating	95% CI	W-L-T
1	demo-claude-code	1.25	[0.79, 1.97]	25-5-0
2	demo-codex-cli	0.60	[0.12, 1.01]	20-10-0
3	demo-aider	-3.28	[-3.59, -3.07]	0-30-0

Cost / Accuracy Tradeoff

WHAT THIS MEANS

The cost/accuracy Pareto frontier plots every harness by what it costs to run a trial against how often it succeeds. A harness is 'Pareto-optimal' when no other harness beats it on both axes at once. A harness that is not Pareto-optimal is strictly dominated: another option exists that is cheaper and at least as accurate, or just as cheap and more accurate.

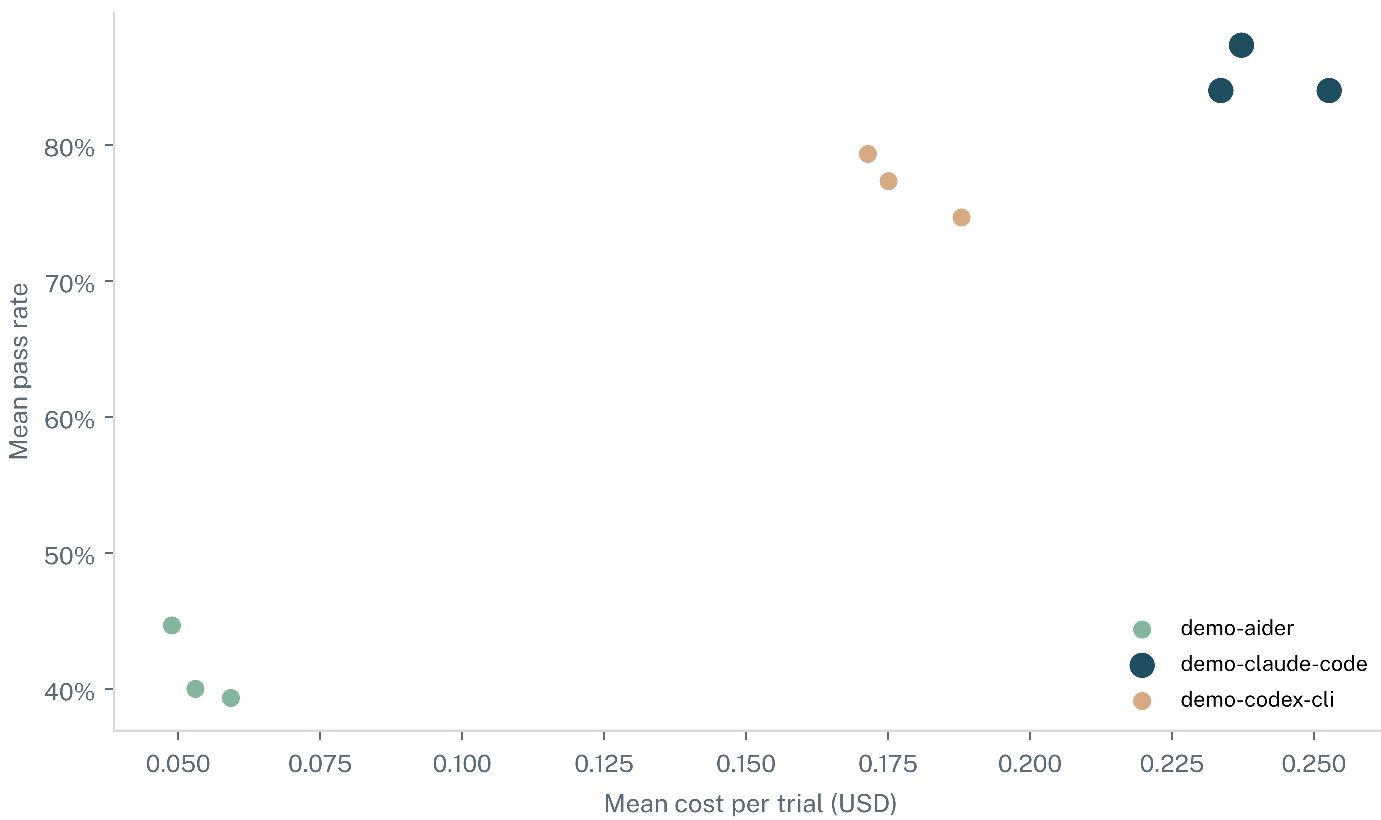

demo-claude-code's standing:

Pareto-optimal — no other harness beats it on both cost and accuracy

Operational Maturity Profile

demo-claude-code's operational shape across five axes, normalized against the other harnesses compared (correctness/speed/cost efficiency are higher-is-better; the last two axes are descriptive, not evaluative).

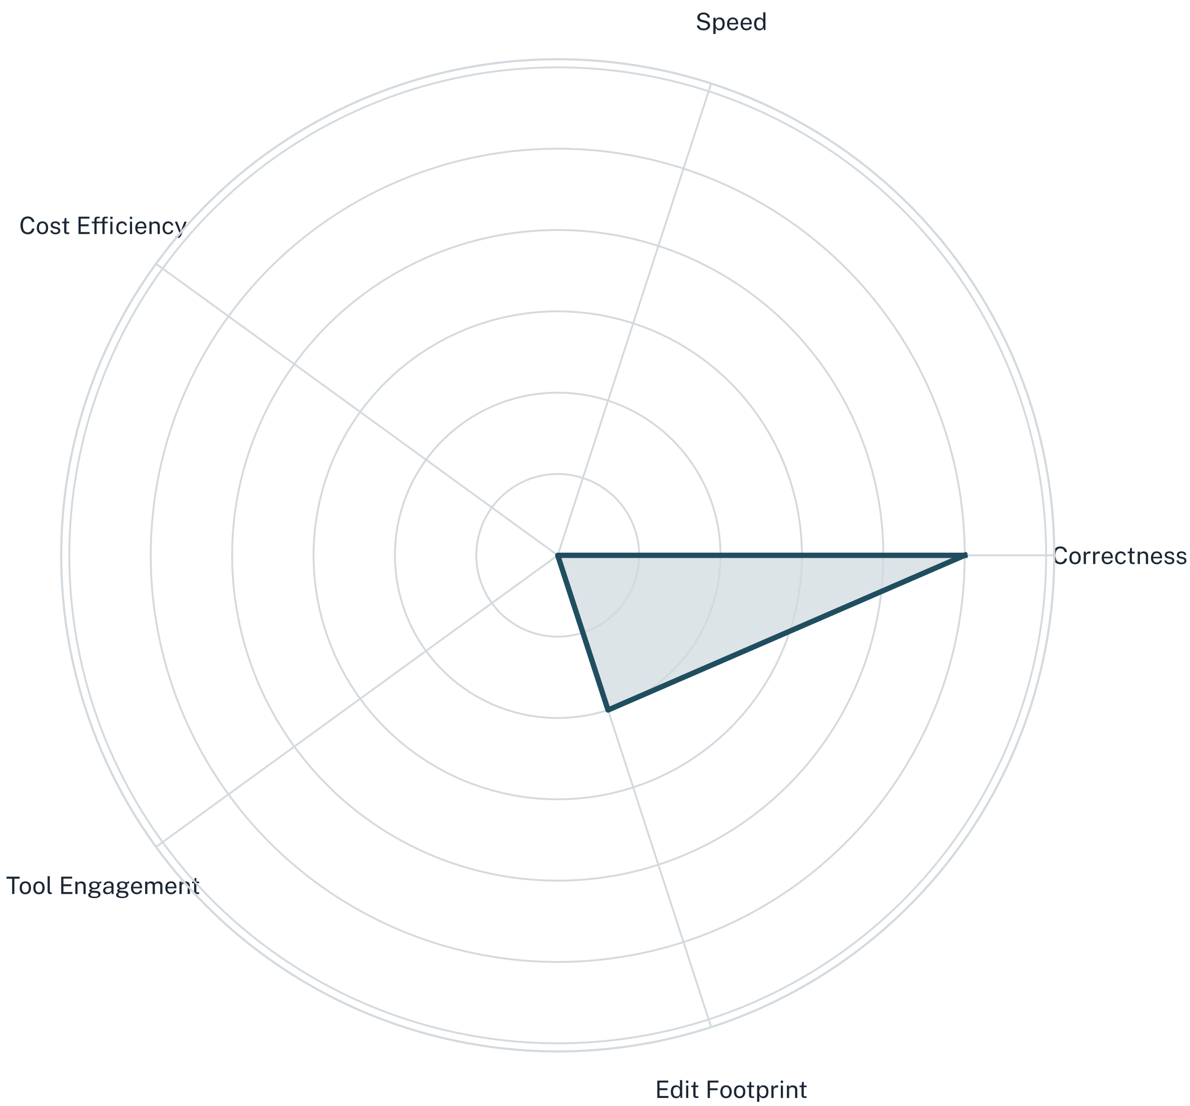

Axis	Normalized score	Raw value
Correctness	1.00	0.851
Speed	0.00	62.9
Cost Efficiency	0.00	0.241
Tool Engagement	0.00	17.9
Edit Footprint	0.40	4.44

Continuous Monitoring

WHAT THIS MEANS

Continuous monitoring tracks a harness's rating over successive submission periods using the Glicko-2 rating system (the same family of methods used for competitive game ratings), so a real capability shift shows up as a trend rather than being lost in a single snapshot. A widening gap from one period to the next is flagged 'watch' or 'drifting' depending on its size relative to the rating's own uncertainty (RD) — this catches a harness quietly regressing (or improving) between formal evaluation rounds.

demo-claude-code (version demo-1.0)

STABLE

latest rating 1658.7 · latest change +10.1 · largest change seen 238.5

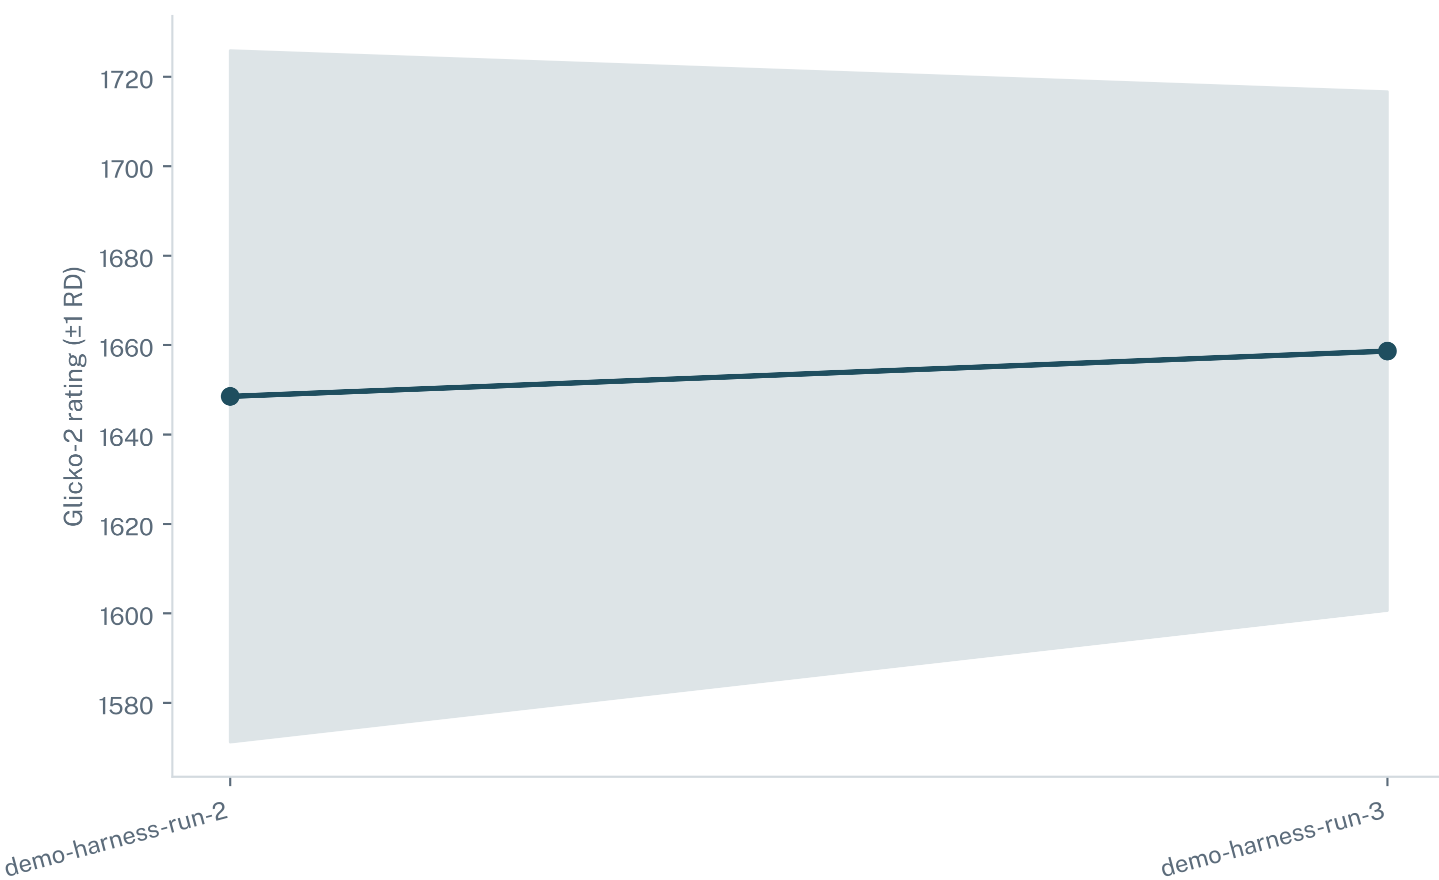

Methodology & Glossary

Bradley-Terry rating

A pairwise-comparison model (the same family behind chess ratings and Chatbot Arena) that estimates each harness's relative strength from how often it beat every other harness on the same (model, task) combination, rather than averaging a raw pass rate that would be distorted by which tasks each harness happened to attempt.

95% confidence interval

The range this report is 95% confident contains the harness's true rating, estimated via bootstrap resampling of the underlying head-to-head comparisons. A narrow interval means high confidence; a wide one means more trials are needed before treating the rank order as settled.

Pareto-optimal

A harness is Pareto-optimal on the cost/accuracy frontier when no other harness beats it on both cost and pass rate simultaneously. See page 4 for this harness's own result.

Glicko-2 continuous monitoring

A rating-with-uncertainty system (rating + RD, rating deviation) tracked across submission periods so a real shift in a harness's performance shows up as a trend, not just a single snapshot. See page 6 for this harness's own drift history.

DATA & METHOD

This report was generated from synthetic demo data for design review, not a real harness evaluation submission. Data shown here should not be used for a real procurement decision.